

9Back : Making Our Plans Great Again

Maxwell Bland

Leon Cheung

Kilian Pfeiffer

University of California, San Diego

Abstract

The following paper describes a set of benchmarks run on the Plan 9 operating system. In particular, it uses the only actively developed branch, 9FRONT "RUN FROM ZONE!" (2018.09.09.0 Release). The goal of this project is to provide researchers and enthusiasts with greater insight into the current state of the operating system's performance and capabilities of interaction with hardware. Through tests of CPU, Memory, Network, and File System operations, we will gain insight on bottlenecks in the system's performance and the interactions between low-level (hardware) and high-level (OS) system components. These performance statistics will be contrasted with subjective experiences of "responsiveness".

1 Introduction

Plan 9 from Bell labs is a distributed operating system which emerged around the 1980's. It built upon the ideas of UNIX, but adopted an ideology that "everything is a file". Although the system was marketed in the 90's, it did not catch on, as prior operating systems had already gained enough of a foothold. Eventually, during the 2000's, Bell Labs ceased development on Plan 9, meaning official development halted. Unofficial¹ development continues on the 9front fork of the codebase, with new support for Wi-Fi, Audio, and everything anyone could ever want or need.

The experiments were performed as a group using a shared codebase and a single machine described in the following section. The measurements were performed via programs written in Plan 9's *special* version of C 99 and Alef, both described in the original Plan 9 paper [todocite](#). The compilers used were the x86 versions of Plan 9's C compiler and Alef compiler, 8c and 8al respectively. The compilers were run with no special optimization settings. Version numbers

are not available. Measurements were performed on a single machine running Plan 9 directly from hardware; given the nature of the Plan 9 system, additional metrics could be established for networked file systems and CPU servers; these measurements were not done for sake of simplicity, and because results under these conditions should be inferable from the results cataloged within this paper.

2 Machine Description

We ran this beautiful operating system of the gods on a Thinkpad T420, the machine of true developers.

```
Processor: model, cycle time, cache sizes (L1, L2,
instruction, data, etc.)
Intel(R) Core(TM) i5-2520M CPU @ 2.50GHz (Sandy Bridge)
cache size 3072 KB
L1$ 128 KiB
L1I$ 64 KiB 2x32 KiB 8-way set associative
L1D$ 64 KiB 2x32 KiB 8-way set associative
write-back
L2$ 512 KiB 2x256 KiB 8-way set associative
write-back
L3$ 3 MiB 2x1.5 MiB 12-way set associative
write-back
cpu family 6
model 42
stepping 7
siblings 4
cores 2
fpu yes
fpu_exception yes
fpu_exception yes
bogomips 4986.98
clflush size 64
cache_alignment 64
Details at https://en.wikichip.org/wiki/intel/core\_i5/i5-520m
```

```
Memory bus
DDR3-1333
```

¹This is debatable. If you adopt an orphan, are they not your official child?

```

i/o-bus-frequency: 666MHz
bus-bandwidth: 10656 MB/s
memory-clock: 166MHz
Column Access Strobe (CAS) latency:

I/O bus
SataIII-speed: 600MB/s

RAM size
8 GB
Disk: capacity, RPM, controller cache size
Samsung SSD 860 EVO 500GB
Capacity: 500GiB
RPM: 550MB/s read, 520 MB/s write
and 98,000 IOPS (Read QD32)
Controller Cache Size: 512MB
Network card speed:
Intel 82579 LM Gigabyte: 1Gb/s
intel Centrino Ultimate-N 6300: 450 Mbps

Operating system (including version/release)
9FRONT "RUN FROM ZONE!" (2018.09.09.0 Release)

```

	Reading nsec	Loop
Hardware Guess	100 nsec	100 nsec
Software Multiplier	10-30	20
Prediction	1-3 usec	2 usec
Average	2.080 usec	2.084 usec
Std Dev.	0.006 usec	0.0034 usec

Figure 1: Measurement Overheads

3 Experiments

For each section, we report the base hardware performance, make a guess as to how much overhead software will add to base hardware performance, combine these two into an overall prediction of performance, and then implement and perform the measurement. In all cases, we run the experiment multiple times, and calculate the standard deviation across measurements. We use the `cycles()` syscall to record the timestamp. Dynamic CPU frequency scaling was disabled for all trials, and all trials were restricted to a single core.

3.1 Measurement Overhead (Reading Time)

The following section reports overheads of reading time. One trial involves looping over 16 `cycles` timing calls, 2^{16} times. We average out for each call of `cycles`, and we do 64 trials altogether.

Since we imagine getting the current cycles clock is a fast operation, so we estimate the hardware performance to be on the scale of nanoseconds, since one clock cycle takes approximately 1 nanoseconds. We think that the `cycles` call should do something similar to reading from the cycle counter, so it may be a bit slower, perhaps on the order of 10ns. [<https://www.7-cpu.com/cpu/SandyBridge.html>]

The software cost of this, however, will be tremendous, since that value must be moved out from a register and potentially to a variable, and a syscall must be made, the result might be around 10 to 30 times the hardware cost.

This would put our predicted cost at around 100 to 300 nanoseconds.

3.2 Measurement Overhead (Loop Overhead)

In order to measure the loop overhead, we took the time at the end of a loop iteration and at the start of a loop iteration, and find the difference between those to measure the loop overhead time. We repeat this 16384 times within each trial, and perform 64 trials. We average over all of these trials.

The cost of doing a single for loop based branch in hardware is also very low, so long as the branch prediction is correct, implying that it is on the order of only a couple of nanoseconds. But there is the need to increment a variable, and check a loop conditional, which adds overhead. 30 to 40 cycles on a 2.4GHz clock is around 16-20 nanoseconds. So this whole check and loop might take 100 ns. [<https://www.7-cpu.com/cpu/SandyBridge.html>]

Software, again, will slow this down, as well as the rest of the operating system. Other system processes may need the processor, data may not be in the cache, and thus the latency might be 20 times this, the same as a reference to the L2 cache, roughly

Thus, our predicted cost is around 2 microseconds.

3.3 Procedure Call Overhead

In measuring the function overhead, we created functions of zero to seven integer arguments, we then proceeded to call each of these functions 20000 times, and found the average time taken over these 20000 calls. Overall, increment overhead was almost non-existent.

The hardware cost of a procedure call using the standard c calling convention should not be that high, as it is almost all done in registers. The caller pushes the arguments to the stack (very little cost), saves the return address, the frame pointer, and calls the function. The hardware cost of this should be more than a for loop, but still reasonable, maybe roughly equitable since there are no comparisons or complex operations. Thus, somewhere on the range of 150ns to 200ns would seem reasonable, maybe even 100ns. [<https://www.7-cpu.com/cpu/SandyBridge.html>]

The software cost of this is probably about the same as the measurement overhead and for loop overhead, since we still may need to context switch to another process, as well as handle interrupts, etc. So I would reason that it would be around 20x the hardware cost.

Hardware Guess	150 nsec	
Software Multiplier Guess	20x	
Prediction	3000 usec	
Num Args	Average	Std Dev.
0	2.110 usec	0.191 usec
1	2.104 usec	1.248 usec
2	2.118 usec	0.346 usec
3	2.113 usec	0.571 usec
4	2.108 usec	0.213 usec
5	2.120 usec	0.915 usec
6	2.111 usec	2.165 usec
7	2.108 usec	0.321 usec

Figure 2: Procedure Call Overheads

	System Call Overhead
Hardware Guess	100 nsec
Software Multiplier	30-40x
Prediction	3-4 usec
Average	4.835 usec
Std Dev.	0.056 usec

Figure 3: Syscall Overhead

This puts our prediction for the cost to be somewhere between 2000 nsec and 4000 nsec. Arguments shouldn't really effect the speed of a function call under our system.

3.4 System Call Overhead

To measure the syscall overhead, we used the `errstr(2)` syscall, in particular, `rerrstr(char*, vlong)`. `reerrstr` reads the error string for the error of a previous syscall, and does not clear the `errstr` buffer. Within each trial, we performed the syscall 16384 times, and averaged over 64 trials.

Data from IBM puts the cost of making a system call on the Pentium processor in 1997/1995 to be 223 cycles, or roughly 1.68 usec for that processor. Our processor is much faster, putting the timing at around 100 nsec.

It is likely that the overhead of a syscall beyond this is pretty high due to the actual operations and implementation of the syscall operation. Thus, the cost could be 30 to 40x the hardware cost.

That puts our estimate at around 3 to 4 usec.

3.5 Task Creation Time

We measured process creation overhead in two different ways, one in which we `rfork`'ed a new process without copying the parent state to the child, and, `fork`, where the child will copy all of the parent's file descriptors, and share the other state as normal. We do these 20000 times for each type of fork and we average over all of these calls. This

	Light Fork	Heavy Fork
Hardware Guess	500 usec	500 usec
Software Multiplier	2x	2x
Prediction	1000 usec	1000 usec
Average	1080 usec	1114 usec
Std Dev	278 usec	344 usec

Figure 4: Fork Overheads

methodology results from the fact that Plan 9 does not have kernel threads.

The hardware costs of a light fork should be much less than a heavy fork, since resources do not need to be shared, but because our processes are minimal, it is likely this difference is negligible. Still, this whole process will take tens of thousands to millions of cycles, meaning thousands of microseconds, maybe 500 usec with a 1,000,000 cycles. [<http://ithare.com/infographics-operation-costs-in-cpu-clock-cycles/>]

The software costs are most likely pretty low since Plan 9's `rfork` mechanism is relatively light weight and similar to the Unix version of `fork`. It might only scale this already tremendous runtime by a factor of two or so.

That puts our estimate at around 1000 usec for both fork operations.

3.6 Context Switch Time

Finally, to measure context switching, we created a pipe in a parent process, and forked off a new child process. We forced a context switch by writing to the pipe in the parent and then waiting for the child to read from the pipe and terminate. We take the time between when the parent goes to sleep and when the parent wakes up again after the child's termination. Then, we subtract the overhead of reading from the pipe and closing the pipe (which happens in the child), and divide the time in half to account for the two context switches. We repeat this measurement 1000 times and use a new pipe each time. To take the time of a context switch, we take the minimum time over all of these trials.

We estimated the base hardware cost of context switching while pinned to a single CPU to be on the order of around 1100 ns based upon benchmarks done on the Intel E5-2620, another Sandy Bridge processor model. [<https://blog.tsuninet.net/2010/11/how-long-does-it-take-to-make-context.html>]

The software cost to a context switch for Plan 9 is going to be huge. Probably on the order of 40x or something along these lines, since the OS needs to deal with scheduling and piping and I/O and virtual memory, etc..

This puts our prediction right around 44 usec.




Figure 5: Ram access time with size of the memory region accessed on x axis in an logarithmic scale and latency in seconds on the y axis

Hardware Cost	1100 nsec
Software Multiplier	40x
Prediction	44 usec
Average	42.048 usec
Std Dev.	9.569 usec
Min	12.491 usec

Figure 6: Context Switch Overheads

3.7 RAM Access Time

To measure ram access time allocated a 1.6 GB array on the heap, and iterated through it with changing stride size to hit L1, L2 cache and finally memory. The stride size gets increased by factor 1.1 for each iteration.

3.8 Discussion

Compare the measured performance with the predicted performance. If they are wildly different, speculate on reasons why. What may be contributing to the overhead? Our hardware estimates were very, very rough. It'd be worthwhile to go back and double check these and discuss with our classmates on the numbers we arrived at.

Evaluate the success of your methodology. How accurate do you think your results are? Our results agree with the Plan 9 paper, which is suprising given that our processor is a bit faster. It will be worth it to discuss with Geoff our results and what we might try looking into going forward. It was hard to ensure we were taking the best approach. We certainly measured roughly what we wanted to measure.

Answer any questions specifically mentioned with the operation. The issue here is that there are no kernel threads in Plan 9, so some of the CPU measurement questions did not apply

4 Conclusion

We should try to correct our idea of hardware cycle counts for typical instructions and discuss the manners in which we chose to benchmark in order to make sure they are the absolute best way to measure these values, or something close to it.

References